



RESULTS FOR QUERY:

Full Report by Use Cases – Expert Transcripts

CREATED

May 19, 2026

PAGES

9

RESEARCH URL

[Click to View In GenSearch](#)

QUERY

ok you gave me the brief report. now give me the full report

TABLE OF CONTENTS

1. Key Takeaways
2. Use Case: High-Level Planning and Project Scoping
3. Use Case: Continuous Integration and Automated Troubleshooting
4. Use Case: Asynchronous Ticket Execution and PR Generation
5. Use Case: Feature Prototyping, Complex Migrations, and Environment Building

Devin's Engineering Applications: Customer Sentiment Across Core Use Cases

Key Takeaways

- **Natural language interpretation stands out as Devin's strongest capability** , with the VP of Engineering at Publicis Sapient noting it is unmatched in understanding engineering intent, actively resolving ambiguities and breaking requirements into actionable tasks . 
- **CI/CD automated troubleshooting delivers measurable ROI** , reducing medium-level error resolution from 2-3 hours to 30-60 minutes, with Devin successfully generating viable PRs for 40-50% of automated alerts at companies like Ramp . 
- **Devin Search enables cross-functional access to engineering knowledge** , with 5-10 non-technical employees at Crossmint using it via Slack to investigate system behaviors, while product managers and designers handle small fixes independently . 
- **Asynchronous ticket execution enables massive task parallelization** , allowing engineers to theoretically assign dozens of tickets simultaneously without hard limits on concurrent sessions, with projections of assigning up to 20 tickets each morning . 
- **PR merge rates reveal significant quality gaps** , with only 35-40% of Devin's 1,500 PRs at Linktree deemed high-quality enough for production, and less than half merged without requiring human intervention . 

- **Complex architectural migrations remain beyond Devin's capabilities** , with experts noting it is not recommended for heavy-lifting tasks like microservices migrations or framework shifts, as the agent frequently goes off the rails and burns through compute credits . 
- **Customer consensus positions Devin as a "junior AI engineer"** that cannot independently complete complicated tasks or massive refactoring, but drastically accelerates medium-complexity features while requiring human engineers for high-level development work . 

Use Case: High-Level Planning and Project Scoping

Engineering organizations are increasingly deploying Devin at the earliest stages of the software development lifecycle to handle high-level technical planning and project scoping. ^[1] Rather than functioning merely as a downstream code generator, the agent is utilized to translate broad concepts into structured engineering tasks. ^[2]

 1 • Expert Call  2 • Expert Call

Natural Language Interpretation and Ambiguity Resolution

A core strength of Devin in the planning phase is its advanced capacity to process unstructured requirements. ^[3] The Vice President of Engineering at Publicis Sapient notes that Devin is currently unmatched in its ability to understand engineering intent and natural language. ^[4] Developers can feed the agent raw user stories, which it then interprets autonomously. ^[5] During this process, Devin actively resolves ambiguities, asks clarifying questions, and breaks the core requirements down into a sequence of actionable planning tasks. ^[6] This allows the tool to act almost as a program manager that structures the exact initiatives required to complete a product development cycle. ^[7]

 3  4  5  6  7 • Expert Call

Feasibility Discussions and High-Level Design

Before manual coding begins, engineering teams are utilizing Devin to validate technical feasibility and architectural approaches. ^[8] At Linktree, engineering leadership has observed a shift in how developers initiate projects, moving toward conversational scoping via the "Ask Devin" user interface. ^[9] This feature functions as a question-and-answer environment where the agent provides technical plans rather than immediate code execution. ^[10] Engineers use this interface to discuss feasibility, steer the high-level design direction, and agree on an overarching approach. ^[11] Once the developer and the agent are perfectly aligned on the architectural plan, the engineer simply clicks a button to transition Devin into the active coding phase. ^[12]

 8  9  10  11  12 • Expert Call

Devin Search and Codebase Indexing

A major component of successful project scoping involves understanding existing architectures, a workflow heavily enhanced by Devin's repository indexing features. ^[13] The Director of AI Engineering at Elementor highlights that Devin leverages an open-source tool called Deep Wiki to provide comprehensive, up-to-date README files for the entire codebase. ^[14] This automated indexing maps out high-level architectures, cron job timings, and system diagrams, allowing technical staff to query exactly how the product currently works before proposing changes. ^[15]

 13  14  15 • Expert Call

This indexing capability has also expanded Devin's scoping utility to non-technical users across the enterprise. ^[16] The Co-founder of Crossmint reports that their company heavily utilizes "Devin Search," an agentic search function that operates directly over their code repository. ^[17] At Crossmint, five to ten non-technical employees, such as support staff, use Devin Search via Slack to investigate system behaviors and answer external client questions.

RESULT FOR QUERY

[18] Similarly, product managers, designers, and data analysts can use the agent to handle small frontend scoping or fixes independently, preventing disruptive context switching for the core engineering team. [19]

16

17

18 • Expert Call

19 • Expert Call

In data engineering contexts, Devin is deployed to map out complex data models and definitions during the planning phase. [20] The Director of Operations & Analytics at Tavily notes that the agent is incredibly effective at crawling through concentric layers of code to determine the exact logic behind specific business metric definitions. [21] This enables teams to validate their architectural assumptions and thoroughly scope out analytical queries before executing massive SQL tasks. [22]

20

21

22 • Expert Call

Limitations in Complex Project Scoping

Despite these robust planning capabilities, experts note distinct limitations in Devin's ability to scope and execute highly complex or context-heavy projects. [23] The Director of Engineering at Linktree points out that Devin struggles to absorb context beyond the raw codebase, severely lacking the business intuition and domain knowledge that human engineers naturally gather from organizational communication. [24] Furthermore, if a planned task is too broad or requires context across more than a thousand files, the agent frequently experiences hallucinations and fails to produce usable outputs. [25]

23 • Expert Call

24

25 • Expert Call

An expert from Anthropic echoed this sentiment regarding complex task execution. [26] While the user experience for assigning tasks is excellent, the underlying foundation models still lack the independent reasoning capabilities of a senior software engineer when navigating highly ambiguous parameters. [27] Consequently, if the model cannot be steered correctly during the initial planning phase, it can burn through compute credits while delivering substandard architectural results. [28]

26

27

28 • Expert Call

Use Case: Continuous Integration and Automated Troubleshooting

Beyond early-stage project scoping, engineering teams are increasingly integrating Devin directly into their continuous integration and continuous deployment (CI/CD) pipelines to function as an autonomous background troubleshooter. [29] In these environments, Devin operates without requiring a human developer to manually assign a ticket. [30] Instead, organizations configure their systems to automatically trigger a Devin session whenever an automated test fails or a system monitor alerts the team to an error. [31] The Director of AI Engineering at Elementor notes that because Devin operates within its own dedicated virtual machine, it can independently spin up a Docker container, run backend services, and execute the application to directly interface with CI pipeline issues. [32] This capability allows the agent to act as a first line of defense, automatically digesting stack traces and attempting to resolve failures before a human engineer ever opens the alert. [33]

29

31 • Expert Call

30 • Expert Call

32 • Expert Call

33 • Expert Call

Autonomous Bug Investigation Mechanics

When a pipeline failure occurs, Devin utilizes its sandbox environment to autonomously investigate the root cause and iterate on potential solutions. [34] A Senior Manager at Scale AI explains that the agent analyzes the

RESULT FOR QUERY

corresponding Jira ticket or error log, understands the relevant codebase, and formulates a plan to edit the code. [35] Crucially, Devin then automatically runs the builds and iterates continuously until the CI/CD tests pass, eventually generating a pull request (PR) for human review. [36] This workflow has become highly valuable for resolving very precise CI/CD failures, flaky background tasks, and basic linter warnings. [37] At Tecton, the VP of Engineering reported that Devin organically began diagnosing test failures—explaining why a test broke and providing the exact fix—which ultimately became one of the agent's most useful unexpected capabilities. [38]

 34  35  36  37 • Expert Call  38 • Expert Call

Custom Configurations: 'Dot Files' and Base Memories

To effectively operate within complex, proprietary CI/CD environments, Devin relies heavily on customizable context configurations known as "dot files" or base memories. [39] A Senior Software Engineer at Ramp details how their teams have mapped out specific instructional files that dictate exactly how the AI should interact with given areas of their codebase. [40] These memories serve as foundational instructions wrapped into the agent's system prompt, allowing developers to preload explicit commands. [41] For example, a memory might instruct Devin that to test a specific API, it must open the command line, run five precise commands in a specific sequence, and parse the resulting output. [42]


 39  40  41  42 • Expert Call

This contextual pre-loading prevents the agent from guessing how to run local environments and dramatically improves its efficacy in self-contained troubleshooting scenarios. [43] Furthermore, the Cofounder at Crossmint observed that Devin continuously updates this contextual knowledge base with every successful PR, learning organizational nuances such as where to look when a specific CI pipeline gets stuck. [44] Over time, users can manually specify additional generic memories for the entire company or tailored memories for specific engineering teams to continuously refine the agent's accuracy. [45]

 43  45 • Expert Call  44 • Expert Call

Success Rates and Return on Investment

Deploying Devin for automated maintenance has yielded measurable time savings for standard, repetitive bug fixes. [46] According to Scale AI, fixing a medium-level CircleCI error might normally take a junior engineer two to three hours to resolve. [47] Devin is frequently able to automate this same process and deliver a solution within 30 to 60 minutes. [48] At Ramp, this automated CI troubleshooting accounts for 20% to 30% of their total Devin usage, successfully transforming close-ended problems into easily reviewable tasks. [49] The expert from Ramp notes that Devin is currently capable of generating a viable PR for roughly 40% to 50% of the automated alerts it tackles, significantly reducing the investigation burden on human developers. [50]

 46  47  48 • Expert Call  49 • Expert Call  50 • Expert Call

Limitations and Human Oversight

Despite these successes, experts warn that the agent is not a complete replacement for human review and struggles significantly when CI failures require broad architectural context. [51] The Director of Engineering at Linktree highlights that less than half of Devin's generated PRs are merged directly into production without requiring a human developer to apply nudges, prompts, or manual code touch-ups. [52] Additionally, if a test failure

is tied to a highly complex PR or requires widespread codebase changes, Devin is prone to entering infinite loops where it spends hours attempting to fix an issue without success. ^[53] Consequently, engineering teams primarily restrict Devin's automated CI/CD access to strictly scoped, self-contained tasks where cause-and-effect relationships are clearly defined. ^[54]

 51 • Expert Call 52 • Expert Call 53 • Expert Call 54 • Expert Call

Use Case: Asynchronous Ticket Execution and PR Generation

While automated maintenance handles immediate pipeline failures, the most profound workflow shift enabled by Cognition's Devin is its capacity for asynchronous ticket execution. ^[55] Unlike IDE-first tools like Cursor that require the developer to be present and actively prompting, Devin operates under a fundamentally different autonomous paradigm. ^[56] Developers or non-technical staff can prompt the agent with a task, allowing it to run independently in a sandboxed environment for minutes or hours while the human focuses on other work. ^[57]

 55 56 57 • Expert Call

Instruction Writing and Ticketing Integrations

This asynchronous capability has led engineering teams to formally integrate Devin into their project management workflows, most notably through platforms like Linear. ^[58] When product managers create feature or bug tickets, teams can simply tag Devin directly within the tracking system. ^[59] The agent reads the ticket's title and description, formulates a multi-step plan, and begins executing the work as if it were triggered via a direct developer interface. ^[60]

 58 59 • Expert Call 60 • Expert Call

For many startups, this is fundamentally shifting the engineering process from writing code to writing instructions. ^[61] A Founding Engineer at Corner Health noted that their team is actively transitioning to a model where developers spend their mornings writing highly detailed Linear specifications for Devin to execute. ^[62] The primary motivation for this workflow is the potential for massive parallelization of tasks. ^[63] Because Devin operates independently, engineers can theoretically assign dozens of tickets simultaneously without any hard limits on the number of concurrent sessions. ^[64] Engineering leaders project that in the near future, developers will assign up to 20 tickets to the agent each morning, returning hours later solely to review and merge the resulting code. ^[65]

 61 62 64 65 • Expert Call 63 • Expert Call

The Synthetic Junior Developer Dynamic

In this capacity, Devin acts as a synthetic junior developer embedded directly within the engineering organization. ^[66] The Co-Founder and CEO of [Cal.com](#) highlighted that Devin functions so convincingly as a team member that non-technical staff have occasionally mistaken the agent for a newly hired human employee. ^[67] However, the scope of work successfully completed by the agent strongly skews toward junior-level output. ^[68] The agent is highly effective at identifying established patterns and consistently applying them across multiple files for repetitive data plumbing tasks, but it struggles with principal-level architectural changes. ^[69]

 68 • Expert Call 67 • Expert Call 69 • Expert Call

Human Oversight and PR Generation Metrics

Because the agent operates asynchronously, human oversight via pull request (PR) reviews remains a critical operational bottleneck. ^[70] At Linktree, the engineering organization generated nearly 1,500 PRs through Devin within a single year, though only about 35% to 40% of those were deemed high-quality enough to merge into production. ^[71] If the output is unpromising, developers frequently discard the PR entirely and switch back to manual coding or copilot tools to complete the work. ^[72]

 70 • Expert Call

 71

 72 • Expert Call

Furthermore, reviewing these automated PRs still requires significant technical intuition to determine if the agent correctly understood the broader system context. ^[73] As the volume of AI-generated code increases, reviewing and testing these submissions has become one of the primary constraints in the software development lifecycle.

^[74]  73 • Expert Call  74 • Expert Call

Customer Satisfaction and Limitations with Complex Tasks

Customer satisfaction with this asynchronous model is highly polarized depending on task complexity. ^[75] For low-priority bugs that might otherwise sit in a backlog for months, Devin provides immense value by independently generating solutions with minimal guidance. ^[76] Conversely, for complex or highly ambiguous features, users often find that the agent goes off the rails and burns through compute credits while delivering substandard results. ^[77]

 75 • Expert Call

 76 • Expert Call

 77 • Expert Call

A Senior Software Engineer at Ramp experienced situations where they had to specify instructions in such granular detail that it would have been faster to simply write the code manually rather than prompt the agent. ^[78] Ultimately, while Devin excels as an independent executor of well-defined, medium-complexity tasks, the underlying foundation models do not yet possess the independent reasoning capabilities required to fully replace a human developer on highly complex initiatives. ^[79]

 78 • Expert Call

 79 • Expert Call

Use Case: Feature Prototyping, Complex Migrations, and Environment Building

Beyond background troubleshooting and basic planning, engineering teams are utilizing Cognition's Devin for active code generation and the prototyping of new product features. The platform is designed to function as an autonomous agent that attempts to replace or heavily augment the developer's traditional coding process. ^[80] To achieve this, Devin fundamentally differs from local coding assistants by operating in its own sandboxed environment. ^[81] When assigned a generative task, the agent spins up a dedicated virtual machine equipped with a browser and a dedicated branch of the organization's codebase. ^[82] This environment-building capability allows developers to hand off a feature request and let the agent work independently for minutes or hours to construct and test a prototype. ^[83]

 80 • Expert Call

 81

 82 • Expert Call

 83 • Expert Call

Spinning Up Endpoints and Leveraging Existing Functions

When configured correctly, Devin has proven highly capable of navigating existing codebases to build functional new features. A Senior Software Engineer at Ramp noted that their team achieves significant success by using

RESULT FOR QUERY

custom "dot files" or base memories to guide the agent. ^[84] By mapping out exactly how the AI should interact with specific areas of the repository, developers can instruct Devin to spin up entirely new endpoints for specific purposes. ^[85]  84  85 • Expert Call

Guided by these contextual memories, the agent excels at codebase exploration, successfully identifying and leveraging pre-existing functions to build out the new endpoints without duplicating logic. ^[86] The Director of AI Engineering at Elementor confirmed that Devin is highly effective for these contained feature developments, reliably handling tasks such as changing application routes, adjusting dates, and modifying user interface components. ^[87] Furthermore, the agent has been described as a breakthrough tool for handling complex data engineering and analytics tasks within these controlled environments. ^[88]  86 • Expert Call  87 • Expert Call  88 • Expert Call

Limitations in Complex Architectural Migrations

Despite successes in localized feature prototyping, experts consistently report severe operational limitations when deploying Devin for complex architectural migrations or large-scale refactoring. While the user experience of tagging the agent and asking it to execute a major migration is seamless, the underlying foundation models currently lack the independent reasoning required for senior-level software architecture. ^[89]  89 • Expert Call

For heavy-lifting tasks like migrating from a monolithic architecture to microservices or shifting entire front-end frameworks, Devin is generally not recommended. ^[90] A Technical Staff Member at Anthropic noted that for framework migrations—such as moving from Ruby to TypeScript across multiple files—human-in-the-loop tools like Claude Code currently perform significantly better than autonomous agents. ^[91]  90  91 • Expert Call

The core issue stems from Devin's inability to maintain stable architectural alignment when confronted with massive context shifts. If assigned a highly complex, multi-file migration, the agent frequently goes off the rails and cannot be easily steered back into the correct direction. ^[92] In these scenarios, the autonomous nature of the tool becomes a liability, as it can spend hours burning through compute credits only to deliver substandard or entirely broken results. ^[93]  92  93 • Expert Call

The "Junior AI Engineer" Reality

Ultimately, the consensus among experts is that Devin's current generative capabilities firmly align with those of a junior software developer. The agent cannot fully or independently take complicated tasks or massive refactoring initiatives to completion. ^[94] While it drastically accelerates the delivery of medium-complexity features and small bugs, human engineers are still absolutely required to execute the complex, high-level development work necessary to maintain competitive digital products. ^[95] Until foundation models achieve higher tiers of independent reasoning, engineering organizations will largely restrict Devin's active coding permissions to self-contained feature generation rather than broad system overhauls. ^[96]  94 • Expert Call  95  96 • Expert Call

Citations

[1] Expert Call • Lovable • 22 Jul 25 • "Vice President at BNY Mellon Sees AI Coding Tools Enhancing Efficiency but Notes Challenges in Privacy and Complexity"

RESULT FOR QUERY

[2] [3] [4] [5] [6] [7] Expert Call • Cognition • 13 Jul 25 • "*Vice President of Engineering at Publicis Sapient Sees Devin as a Key AI Tool for Enhancing Developer Productivity and Efficiency*"

[8] [9] [10] [11] [12] [58] [59] [63] [69] [71] [72] Expert Call • Cognition • 24 Sep 25 • "*Director of Engineering at Linktree Believes Devin Cognition Leads in AI Agent Productivity and Training Benefits*"

[13] [14] [15] [19] [32] [87] [94] Expert Call • Cursor Oy • 17 Feb 26 • "*Director of AI Engineering at Elementor Believes AI Tools Enhance Workflow but Lack Concrete ROI Metrics*"

[16] [17] [18] Expert Call • Cognition • 24 Sep 25 • "*Cofounder at Crossmint Sees Devin Outperforming Cursor in AI Tool Evaluation for Start-Up*"

[20] [21] [22] [88] Expert Call • Tavily • 11 Feb 26 • "*Director of Operations & Analytics at Lindy Sees Limited Moats for Independent Web Search Providers Amid AI Advancements*"

[23] [26] [27] [28] [77] [81] [82] [89] [90] [91] [92] [93] Expert Call • Cognition • 24 Sep 25 • "*Technical Staff Member at Anthropic Is Optimistic About AI Advancements in Coding Tools by Year-End*"

[24] [25] Expert Call • Cognition • 01 Oct 25 • "*Director of Engineering at Linktree Prefers Investing in Devin for AI Tool Adoption Over Cursor*"

[29] [31] [49] [54] [75] [78] Expert Call • Cognition • 16 Jul 25 • "*Senior Software Engineer at Ramp Believes Cursor and Claude Code Will Coexist, Catering to Different User Preferences*"

[30] [39] [40] [41] [42] [43] [45] [50] [51] [84] [85] [86] Expert Call • Cognition • 10 Jul 25 • "*Senior Software Engineer at Ramp Believes Devin Lags Behind Claude Code and Cursor in Integration and Productivity*"

[33] [34] [35] [36] [37] [46] [47] [48] [53] Expert Call • Windsurf • 15 Sep 25 • "*Senior Manager at Scale AI Sees Growth in AI Tools Like Cursor and Devin, Highlights Advantages Over GitHub Copilot*"

[38] Expert Call • Coderabbit • 21 Oct 25 • "*VP of Engineering at Tecton Sees AI Tools Enhancing Efficiency and Focus in Code Review Process*"

[44] [66] [68] Expert Call • Cognition • 29 Sep 25 • "*Cofounder at Crossmint Believes Devin's Advanced Features and ROI Make It a Valuable Tool for Engineers*"

[52] Expert Call • Cognition • 18 Sep 25 • "*Director of Engineering at Linktree Sees AI Tools Enhancing Productivity but Values Top Engineers for Complex Tasks*"

[55] [56] [57] [83] Expert Call • Cursor Oy • 12 Apr 26 • "*Principal Engineer at Cognite Believes AI-Powered Tools Like Cursor Are Surpassing GitHub Copilot in Developer Workflows*"

[60] [61] [62] [64] [65] Expert Call • Cognition • 13 Jul 25 • "*Founding Engineer at Corner Health Is Optimistic About Devin's Role in Autonomous AI Software Development*"

[67] Expert Call • Replit • 21 Mar 26 • "*Co-Founder & CEO at Cal.com, Inc Believes Replit Lacks Product-Market Fit for Technical Teams, Praises Devin and Claude Code for Enterprise Use*"

[70] Expert Call • Twilio Inc • 01 Aug 25 • "*Tech Lead at Sierra.ai Believes Twilio and NoSQL Databases Offer Stability and Flexibility in Evolving AI Landscape*"

[73] [76] Expert Call • Graphite.dev • 13 Aug 25 • "*Director of Engineering at Linktree Sees Devin Tool Enhancements Boosting Code Review Efficiency*"

[74] [79] [80] [95] [96] Expert Call • Github • 13 Apr 26 • "*VP of AI at HackerOne Sees AI Tools Transforming Software Development but Notes Limitations in Autonomy*"

Legal Notices

This content was generated using AI and may contain errors, ambiguities, or outdated information.

AlphaSense does not guarantee the accuracy or completeness of this content.

Users assume full responsibility for any use of or reliance upon any content contained in the service. Users should independently validate all information before using or relying on it.

This content does not reflect the views or opinions of AlphaSense or any of its personnel and does not constitute legal, investment, or other professional advice.

This content is proprietary and protected under copyright and other intellectual property laws. No part of this content may be used in connection with any language models or generative AI systems outside of AlphaSense. Any unauthorized use, including but not limited to the training or fine-tuning of any language models or other AI models, is strictly prohibited.